

# A Distributed and Parallel Asynchronous Unite and Conquer Method to Solve Large Scale Non-Hermitian Linear Systems

Xinzhe WU<sup>1,2</sup>    Serge G. Petiton<sup>1,2</sup>

<sup>1</sup>Maison de la Simulation/CNRS, Gif-sur-Yvette, 91191, France

<sup>2</sup>CRIStAL, University of Lille 1, Science and Technology

January 29, 2018

HPC Asia 2018, Tokyo, Japan



# Outline

- 1 Introduction, toward extreme computing
- 2 Asynchronous Unite and Conquer GMRES/LS-ERAM (UCGLE) method
- 3 Experimentations, evaluation and analysis
- 4 Conclusion and Perspectives

## Krylov Subspace

$$K_m = \text{span}\{r_0, Ar_0, \dots, A^{m-1}r_0\}$$

### Different Krylov Methods:

- ① Resolution of linear systems
  - ⊖ GMRES
  - ⊖ CG
  - ⊖ BiCG, etc.
- ② Resolution of eigenvalue problems
  - ⊖ ERAM
  - ⊖ IRAM, etc.

# Future Parallel Programming Trends

## Future Programming Trends:

- ➊ Highly hierarchical architectures
  - ⊖ Computing
  - ⊖ Memory
- ➋ Increasing levels and degree of parallelism
- ➌ Heterogeneity
  - ⊖ Computing
  - ⊖ Memory
  - ⊖ Scalability
- ➍ Requirement of parallel programming
  - ⊖ Multi-grain
  - ⊖ Multi-level memory
  - ⊖ Reducing synchronizations and promoting asynchronicity
  - ⊖ Multi-level scheduling strategies

## Toward extreme computing, some correlated goals

- ⊖ Minimize the global computing time
- ⊖ Accelerate the convergence
- ⊖ Minimize the number of communications
- ⊖ Minimize the number of longer size scalar products and reductions
- ⊖ Minimize the memory space, cache optimization
- ⊖ Select the best sparse matrix compressed format
- ⊖ Mixed arithmetic
- ⊖ Minimize energy consumption
- ⊖ Fault tolerance, resilience

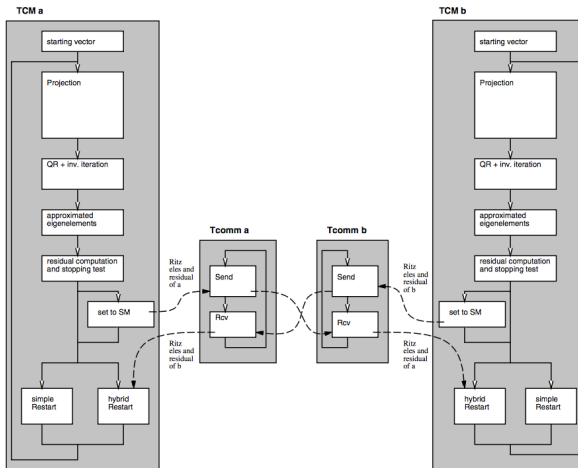
## Toward extreme computing, some correlated goals

- ⊖ Minimize the global computing time
  - ⊖ Accelerate the convergence
  - ⊖ Minimize the number of communications
  - ⊖ Minimize the number of longer size scalar products and reductions
  - ⊖ Minimize the memory space, cache optimization
  - ⊖ Select the best sparse matrix compressed format
  - ⊖ Mixed arithmetic
  - ⊖ Minimize energy consumption
  - ⊖ Fault tolerance, resilience
- Preconditioning
- Unite and Conquer
- 
- ```
graph LR; A[Minimize the global computing time] --> B[Preconditioning]; C[Accelerate the convergence] --> B; D[Minimize the number of longer size scalar products and reductions] --> E[Unite and Conquer]; F[Minimize energy consumption] --> E; G[Fault tolerance, resilience] --> E;
```

# Unite and Conquer Approach

Unite and conquer approach: improving the convergence using other iterative methods[Emad, Nahid and Petiton, Serge, 2016].

**Figure:** Multiple Explicitly Restarted Arnoldi Method (MERAM) [Nahid Emad et al, 2005].



# Outline

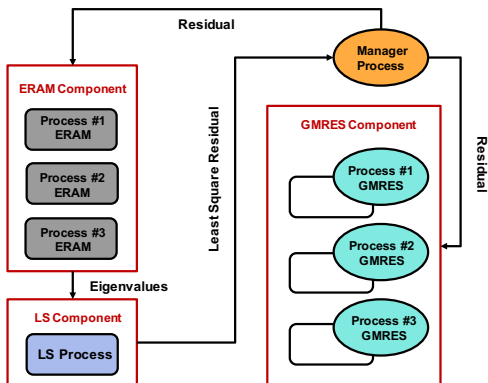
- 1 Introduction, toward extreme computing
- 2 Asynchronous Unite and Conquer GMRES/LS-ERAM (UCGLE) method**
- 3 Experimentations, evaluation and analysis
- 4 Conclusion and Perspectives



# UCGLE Method Implementation

**UCGLE** method is proposed to solve the non-Hermitian linear systems based on the work of this article [Essai, Azeddine and Bergère, Guy and Petiton, Serge G, 1999].

**Figure:** Workflow of UCGLE method





# Least Squares Method

Polynomial preconditioner iterates:  $x_n = x_0 + P_n(A)r_0 \rightarrow r_n = R_n(A)r_0$  with  $R_n(\lambda) = 1 - \lambda P_n(\lambda)$ .

The purpose is to find a kind of polynomial  $P_n$  which can minimize  $R_n(A)r_0$ . For more details of this method, see the article [Youssef Saad, 1987].


**Figure:** Eigenvalues, convex hull and ellipse



# Least Squares Method

## Least Squares method residual

$$r = (R_k(A))^t r_0 = \sum_{i=1}^m \rho((R_k)(\lambda_i)^t) u_i + \sum_{i=m+1}^n \rho((R_k)(\lambda_i)^t) u_i$$

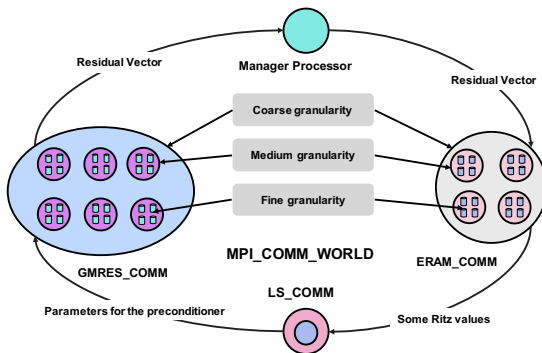


**Figure:** An example of UCGLE for convergence.

# Software engine, orchestration of UCGLE

All three computation components are implemented using the scientific libraries PETSc and SLEPc, based on the work of Pierre-Yves Aquilanti during his thesis at University of Lille 1 [Pierre-Yves Aquilanti, 2011].

**Figure:** Asynchronous Communication and Parallelism of UCGLE method



# Components Implementation

The method implementation is based on the work of Pierre-Yves Aquilanti during his thesis at University of Lille 1 [Pierre-Yves Aquilanti, 2011].

## GMRES Component

The GMRES component is well implemented by the PETSc library.

## Arnoldi Component

The Arnoldi component is implemented by the SLEPc library to calculate the eigenvalues of the matrix operator  $A$ .

## LS Component

Using the Cholesky algorithm, which is provided by PETSc as a preconditioner, but can be used without problem as a factorization method correctly.

# Important Parameters

There are large number of parameters in UCGLE for the users to select and autotune in order to get the best performance.

## I. GMRES Component

- \*  $m_g$ : GMRES Krylov Subspace size
- \*  $\epsilon_g$ : absolute tolerance used for the GMRES convergence test
- \*  $P_g$ : GMRES processors number
- \*  $s_{use}$ : number of times that polynomial applied before taking account into the new eigenvalues
- \*  $L$ : number of GMRES restarts before each time LS preconditioning

## II. Arnoldi Component

- \*  $m_a$ : Arnoldi Krylov subspace size
- \*  $r$ : number of eigenvalues required
- \*  $\epsilon_a$ : convergence tolerance
- \*  $P_a$ : Arnoldi processors number

## III. LS Component


- \*  $d$ : Least Squares polynomial degree

# Outline

- 1 Introduction, toward extreme computing
- 2 Asynchronous Unite and Conquer GMRES/LS-ERAM (UCGLE) method
- 3 Experimentations, evaluation and analysis**
- 4 Conclusion and Perspectives

# Test Matrices

All the following results come from this article [Xinzhe WU and Serge G. Petiton, 2017].



**Figure:** Two strategies of large and sparse matrix generator

**Table:** Test matrices information

| Matrix Name     | n                   | nnz                | Matrix Type   |
|-----------------|---------------------|--------------------|---------------|
| <i>matLine</i>  | $1.8 \times 10^7$   | $2.9 \times 10^7$  | non-Symmetric |
| <i>matBlock</i> | $1.8 \times 10^7$   | $1.9 \times 10^8$  | non-Symmetric |
| <i>MEG1</i>     | $1.024 \times 10^7$ | $7.27 \times 10^9$ | non-Hermitian |
| <i>MEG2</i>     | $5.1 \times 10^6$   | $3.64 \times 10^9$ | non-Hermitian |



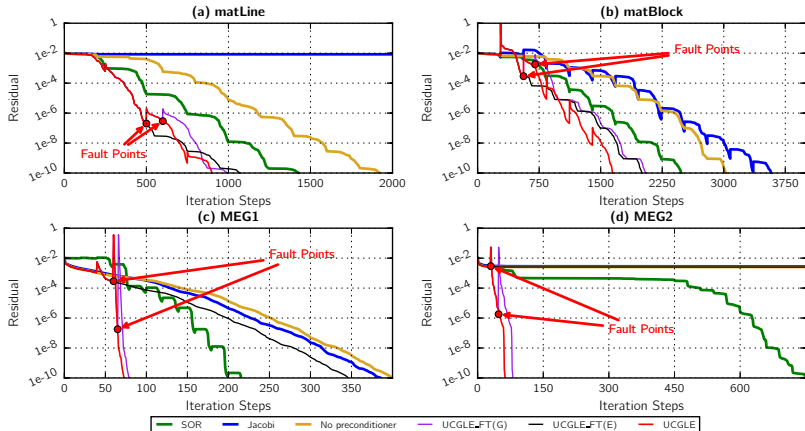
# Experimental Hardware

Experiments on the ROMEO supercomputer in Reims (Champagne, France). ROMEO has 130 nodes. each node has 2 CPU with 8 cores and 2 GPUs. The node specification is given as following:

**Table:** Node Specifications of the cluster ROMEO

|              |                                                     |
|--------------|-----------------------------------------------------|
| Nodes Number | BullX R421 $\times$ 130                             |
| Mother Board | SuperMicro X9DRG-QF                                 |
| CPU          | Intel Ivy Bridge 8 cores 2,6 GHz $\times$ 2 sockets |
| Memory       | DDR 32GB                                            |
| GPU          | NVIDIA Tesla K20X $\times$ 2                        |
| Memory       | GDDR5 6 GB / GPU                                    |

# Convergence and Fault Tolerance Evaluation



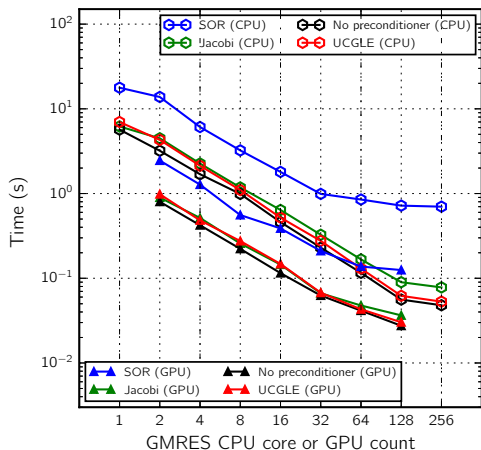
**Figure:** Convergence comparison of *matLine*, *matBlock*, *MEG1* and *MEG2* by UCGLE, classic GMRES, Jacobi preconditioned GMRES, SOR preconditioned GMRES, UCGLE\_FT(G) and UCGLE\_FT(E); X-axis refers to the iteration step for each method; Y-axis refers to the residual, a base 10 logarithmic scale is used for Y-axis.

# Summary of Iteration Number for Convergence

**Table:** Summary of iteration number for convergence of 4 test matrices using SOR, Jacobi, non preconditioned GMRES,UCGLE\_FT(G),UCGLE\_FT(G) and UCGLE: red  $\times$  in the table presents this solving procedure cannot converge to accurate solution (here absolute residual tolerance  $1 \times 10^{-10}$  for GMRES convergence test) in acceptable iteration number (20000 here).


| Matrix Name     | SOR  | Jacobi   | No preconditioner | UCGLE_FT(G) | UCGLE_FT(G) | UCGLE |
|-----------------|------|----------|-------------------|-------------|-------------|-------|
| <i>matLine</i>  | 1430 | $\times$ | 1924              | 995         | 1073        | 900   |
| <i>matBlock</i> | 2481 | 3579     | 3027              | 2048        | 2005        | 1646  |
| <i>MEG1</i>     | 217  | 386      | 400               | 81          | 347         | 74    |
| <i>MEG2</i>     | 750  | $\times$ | $\times$          | 82          | $\times$    | 64    |

# Strong Scalability Results



**Figure:** Strong scalability test of solve time per iteration for UCGLE, GMRES without preconditioner, Jacobi and SOR preconditioned GMRES using matrix *MEG1* on CPU and GPU; X-axis refers respectively to CPU cores of GMRES from 1 to 256 and GPU number of GMRES from 2 to 128; Y-axis refers to the average execution time per iteration. A base 2 logarithmic scale is used for X-axis, and a base 10 logarithmic scale is used for Y-axis.

# Performance Evaluation



**Figure:** Performance comparison of solve time per iteration for UCGLE, GMRES without preconditioner, Jacobi and SOR preconditioned GMRES using matrix *MEG1* on CPU and GPU; X-axis refers respectively to the total CPU cores number or GPU number for these four methods; Y-axis refers to the average execution time per iteration. A base 2 logarithmic scale is used for X-axis, and a base 10 logarithmic scale is used for Y-axis.

# Outline

- 1 Introduction, toward extreme computing
- 2 Asynchronous Unite and Conquer GMRES/LS-ERAM (UCGLE) method
- 3 Experimentations, evaluation and analysis
- 4 Conclusion and Perspectives

# Conclusion and Perspectives

Then

- 1 UCGLE is an asynchronous preconditioned method, minimizing communications, fault tolerant, and allowing efficient GMRES/LS computation for other systems with different right hand sides;
- 2 Several other preconditioners "may be used" (FGMRES) between LS polynomial "accelerations";
- 3 A lot of parameters have to be analyzed : smart-tuning at runtime, learning, toward intelligent linear algebra;
- 4 We have to experiment with very large matrices and on world larger supercomputers to evaluate the impact of large latency for reduction;
- 5 Adapted programming paradigms have to be used for such asynchronous multi-granularity distributed and parallel computing (YML-XMP/YML-XACC, ...);
- 6 SMG2S: generation of Non-Hermitian matrices, including some generate with a given spectrum (soon proposed on line)

# References



Nahid Emad and Serge Petiton (2016)

Unite and conquer approach for high scale numerical computing  
*Journal of Computational Science*, 5 – 14.



Nahid Emad and Serge Petiton and Guy Edjlali (2005)

Multiple explicitly restarted Arnoldi method for solving large eigenproblems  
*SIAM Journal on Scientific Computing*, 253 – 277.



Xinzhe WU, Serge G. Petiton (2018)

A Distributed and Parallel Asynchronous Unite and Conquer Method to Solve Large Scale Non-Hermitian Linear Systems.  
*International Conference on High Performance Computing in Asia-Pacific Region*, accepted.



Azeddine Essai, Guy Bergère and Serge G. Petiton (1999)

Heterogeneous Parallel Hybrid GMRES/LS-Arnoldi Method.  
*PPSC*, 1999.



Youssef Saad (1987)

Least squares polynomials in the complex plane and their use for solving nonsymmetric linear systems  
*SIAM Journal on Numerical Analysis*, 155 – 169.



Pierre-Yves Aquilanti (2011)

Methodes de Krylov reparties et massivement paralleles pour la resolution de problemes de Geoscience sur machines heterogenes dépassant le petaflop.  
*PhD dissertation, Université de Lille*.



Thank you for your attentions!

Questions?